

# **COMP3411 / COMP9414: Artificial Intelligence**

## **11a. Course Review**

# Assessment

---

Assessable components of the course:

Assignment 1	12%
Assignment 2	12%
Assignment 3	16%
Written Exam	60%

# Topics Covered

---

- AI, Tasks, Agents & Prolog
  - ▶ What is AI?
  - ▶ Classifying Tasks
  - ▶ Agent Types
  - ▶ Prolog Programming
- Solving Problems by Search
  - ▶ Path Search
  - ▶ Heuristic Path Search
  - ▶ Games
  - ▶ Constraint Satisfaction
- Logic & Uncertainty
  - ▶ Logical Agents
  - ▶ First Order Logic
  - ▶ Uncertainty
- Learning
  - ▶ Learning and Decision Trees
  - ▶ Perceptrons & Neural Networks
  - ▶ Reinforcement Learning

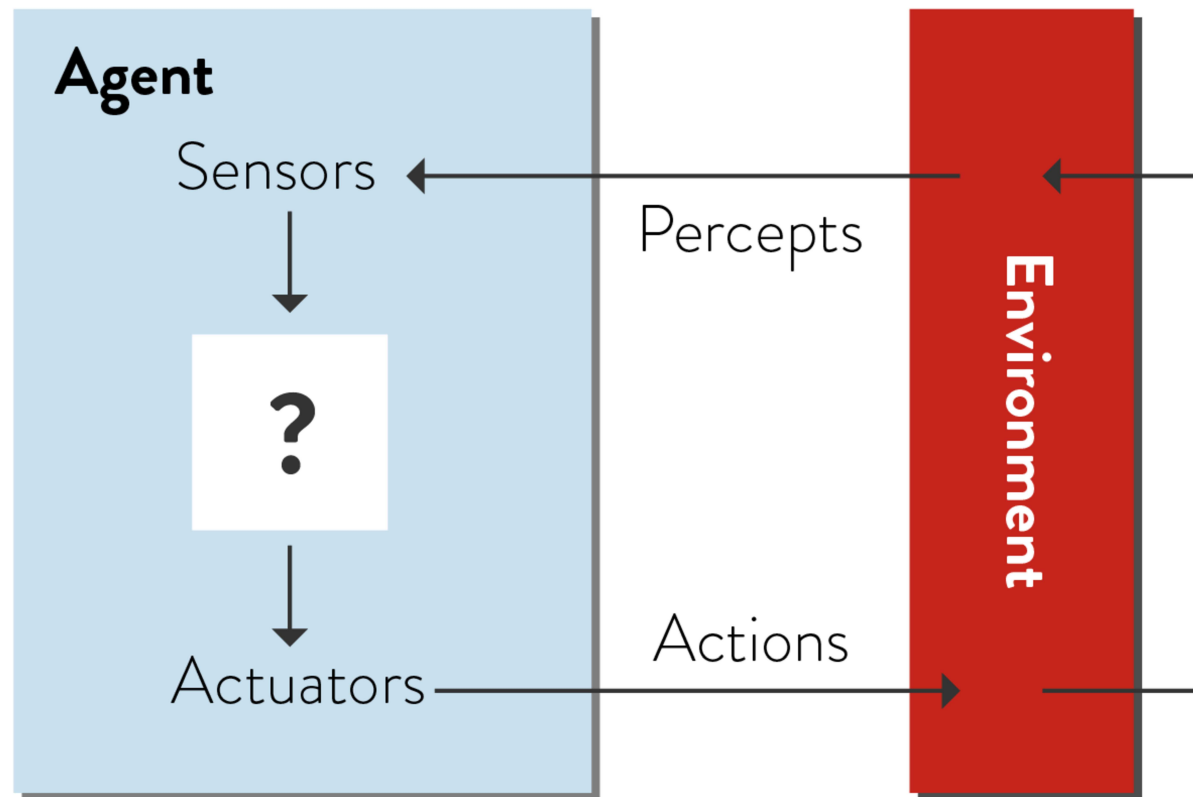
# Not Examinable

---

- Robot Motion Planning
- Evolutionary Robotics
- Computer Vision
- Deep Learning

# Agent Model

---



# Environment types

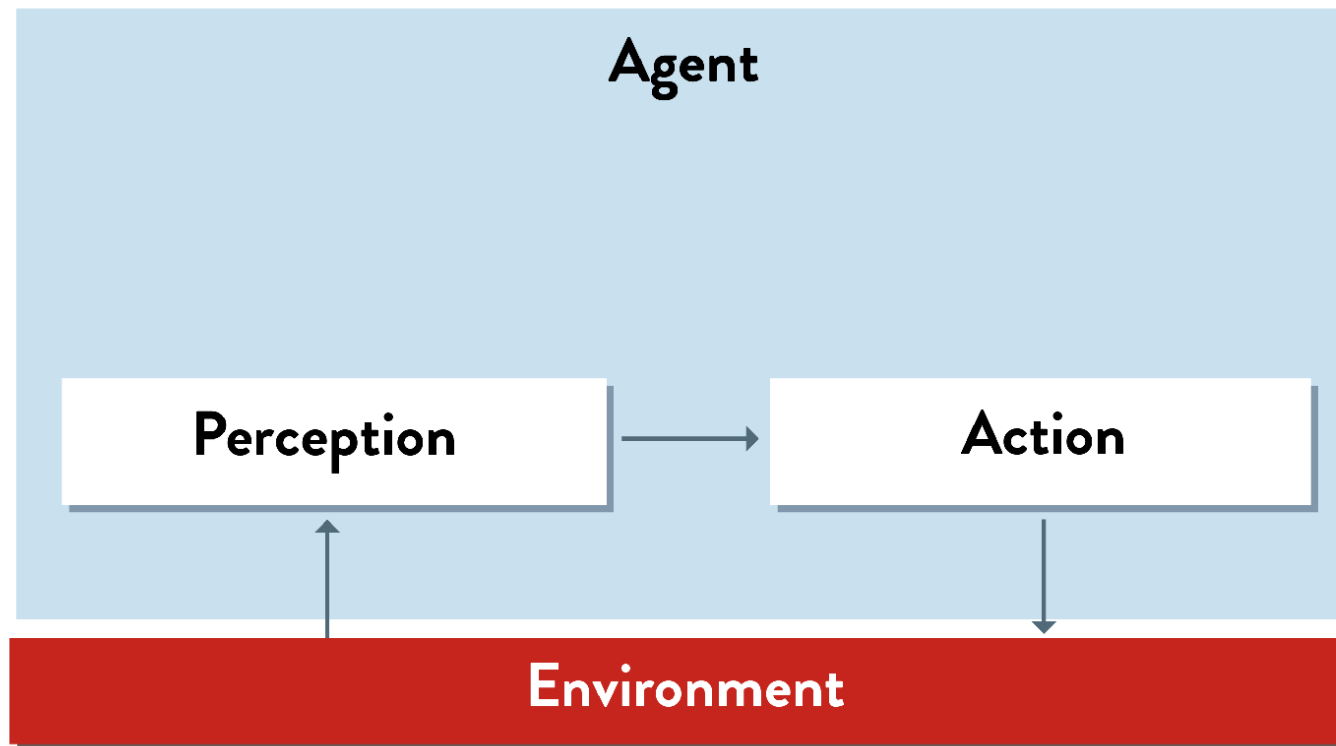
---

We can classify environments as:

- Simulated vs. Situated or Embodied
- Static vs. Dynamic
- Discrete vs. Continuous
- Fully Observable vs. Partially Observable
- Deterministic vs. Stochastic
- Episodic vs. Sequential
- Known vs. Unknown
- Single-Agent vs. Multi-Agent


# Reactive Agent

---



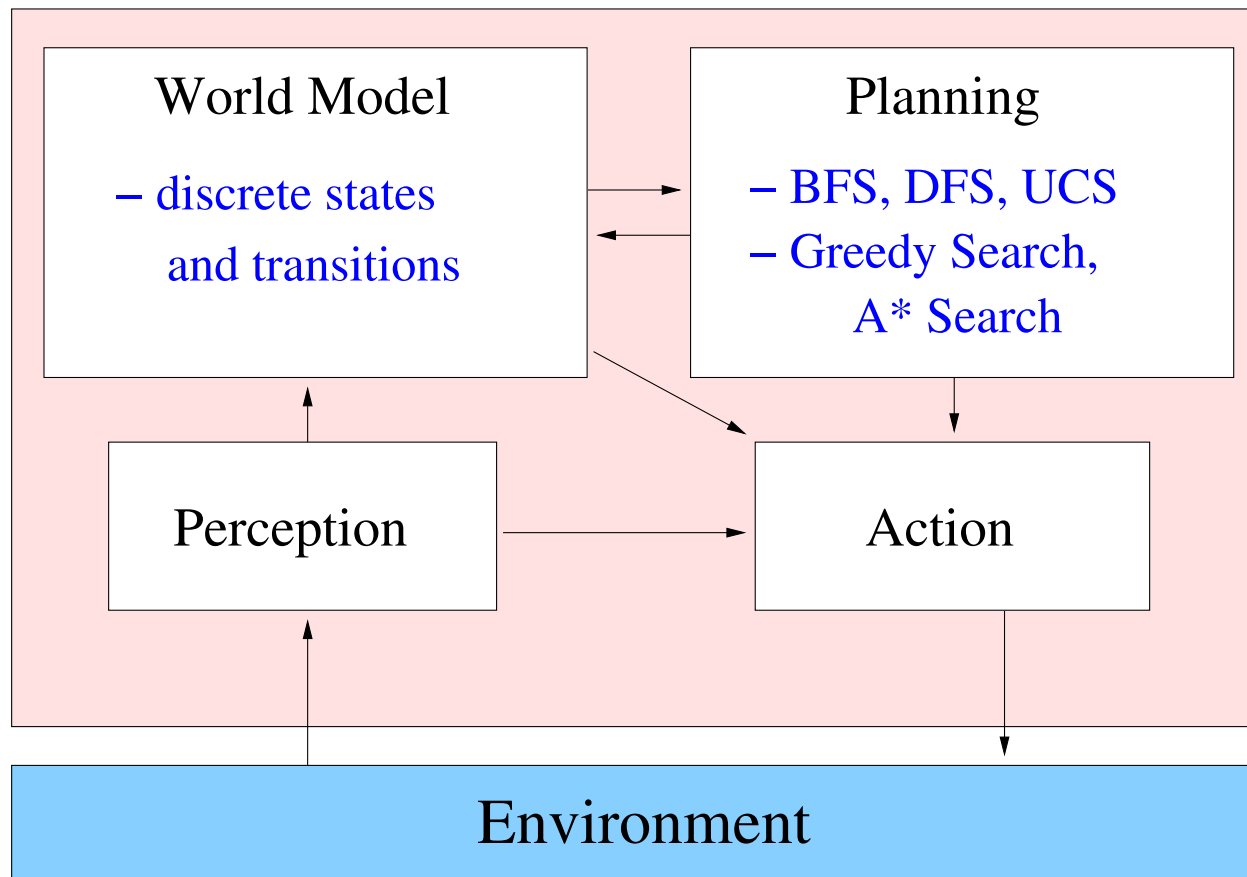
# Braitenberg Vehicles

---





# Path Search Agent



# Path Search Algorithms

---

General Search algorithm:

- add initial state to queue
- repeat:
  - ▶ take node from front of queue
  - ▶ test if it is a goal state; if so, terminate
  - ▶ “expand” it, i.e. generate successor nodes and add them to the queue

Search strategies are distinguished by the order in which new nodes are added to the queue of nodes awaiting expansion.

# Search Strategies

---

- BFS and DFS treat all new nodes the same way:
  - ▶ BFS      add all new nodes to the **back** of the queue
  - ▶ DFS      add all new nodes to the **front** of the queue
- (Seemingly) **Best First Search** uses an evaluation function  $f()$  to order the nodes in the queue; we have seen one example of this:
  - ▶ UCS       $f(n) = \text{cost } g(n) \text{ of path from root to node } n$
- **Informed** or **Heuristic** search strategies incorporate into  $f()$  an estimate of distance to goal
  - ▶ Greedy Search     $f(n) = \text{estimate } h(n) \text{ of cost from node } n \text{ to goal}$
  - ▶ A\* Search         $f(n) = g(n) + h(n)$

# Complexity Results for Uninformed Search

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Time	$O(b^d)$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^k)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bk)$	$O(bd)$
Complete?	Yes <sup>1</sup>	Yes <sup>2</sup>	No	No	Yes <sup>1</sup>
Optimal ?	Yes <sup>3</sup>	Yes	No	No	Yes <sup>3</sup>

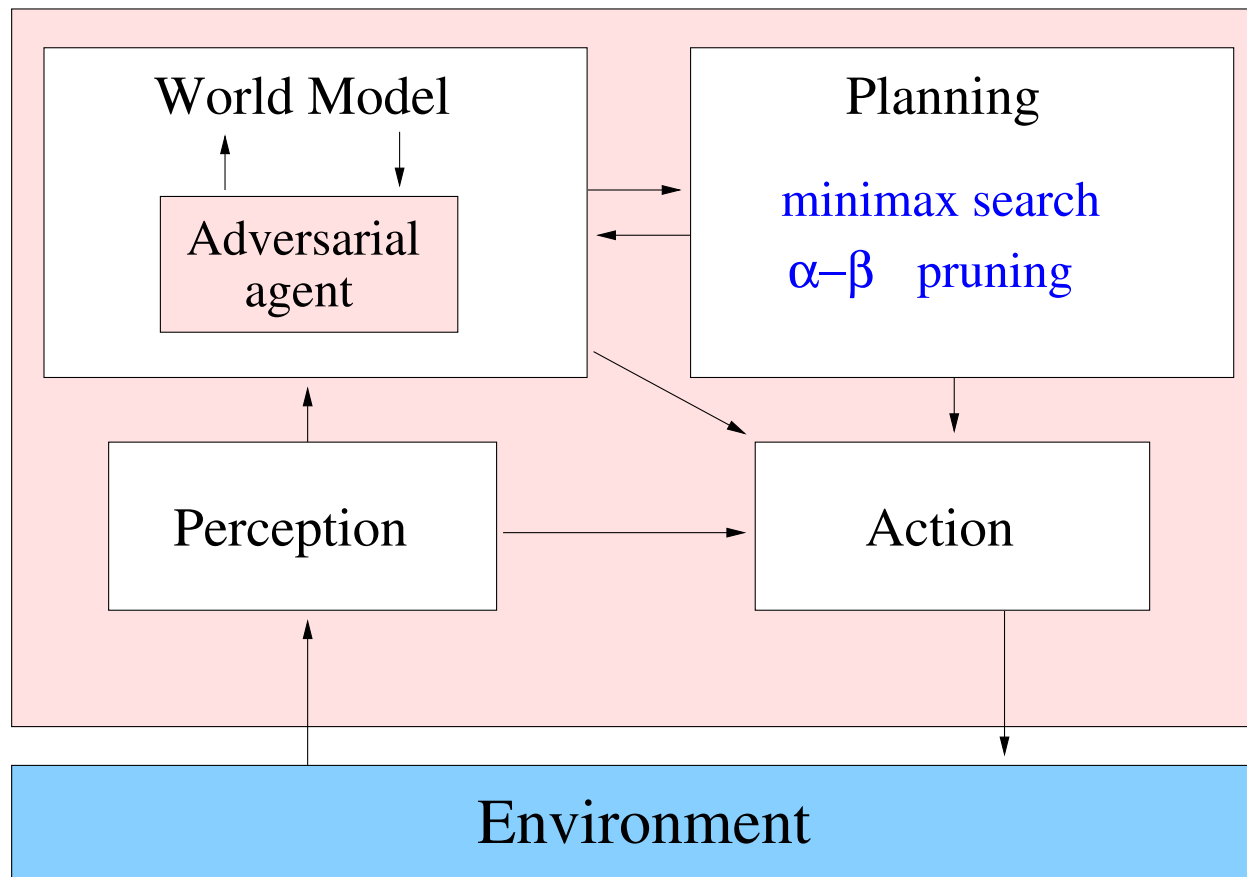
$b$  = branching factor,  $d$  = depth of the shallowest solution,  
 $m$  = maximum depth of the search tree,  $l$  = depth limit.

1 = complete if  $b$  is finite.


2 = complete if  $b$  is finite and step costs  $\geq \epsilon$  with  $\epsilon > 0$ .

3 = optimal if actions all have the same cost.


# Game Search Agent



# Minimax Search




# $\alpha$ - $\beta$ pruning



# Constraint Satisfaction Problems

---

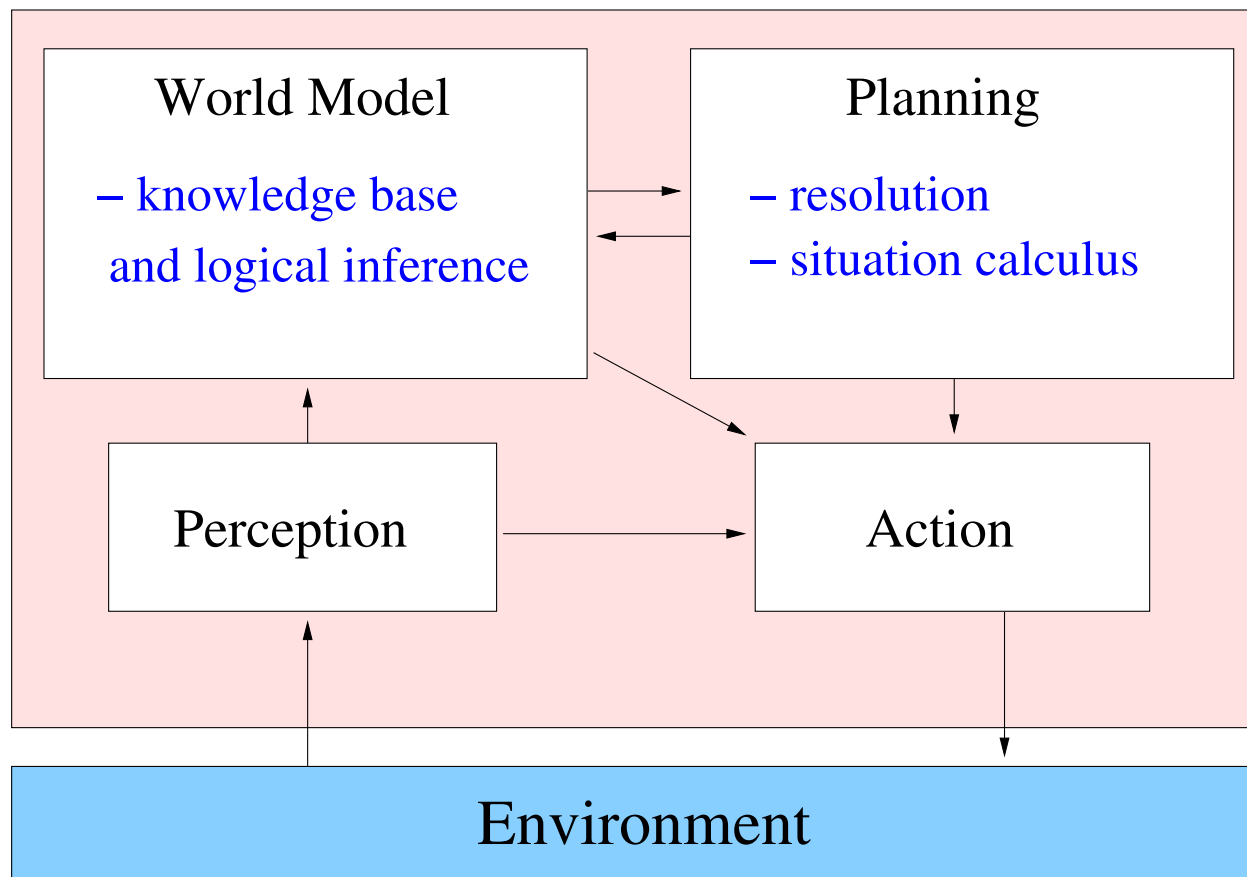


- backtracking search
- enhancements to backtracking search
- local search
  - ▶ hill climbing
  - ▶ simulated annealing



# Logical Agent

---



# Propositional Logic

---

A sentence is **valid** if it is true in **all** models,

e.g. TRUE,  $A \vee \neg A$ ,  $A \Rightarrow A$ ,  $(A \wedge (A \Rightarrow B)) \Rightarrow B$

Validity is connected to inference via the **Deduction Theorem**:

$KB \models \alpha$  if and only if  $(KB \Rightarrow \alpha)$  is valid

A sentence is **satisfiable** if it is true in **some** model

e.g.  $A \vee B$ ,  $C$

A sentence is **unsatisfiable** if it is true in **no** models

e.g.  $A \wedge \neg A$

Satisfiability is connected to inference via the following:

$KB \models \alpha$  if and only if  $(KB \wedge \neg \alpha)$  is unsatisfiable

i.e. prove  $\alpha$  by *reductio ad absurdum*

# Truth Tables

---

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$
F	F	T	F	F	T
F	T	T	F	T	T
T	F	F	F	T	F
T	T	F	T	T	T

# Resolution

## Conjunctive Normal Form (CNF – universal)

conjunction of disjunctions of literals  
clauses

e.g.  $(A \vee \neg B) \wedge (B \vee \neg C \vee \neg D)$

**Resolution** inference rule (for CNF): complete for propositional logic

$$\frac{\ell_1 \vee \cdots \vee \ell_k, \quad m_1 \vee \cdots \vee m_n}{\ell_1 \vee \cdots \vee \ell_{i-1} \vee \ell_{i+1} \vee \cdots \vee \ell_k \vee m_1 \vee \cdots \vee m_{j-1} \vee m_{j+1} \vee \cdots \vee m_n}$$

where  $\ell_i$  and  $m_j$  are complementary literals. e.g.

$$\frac{P_{1,3} \vee P_{2,2}, \quad \neg P_{2,2}}{P_{1,3}}$$

Resolution is sound and complete for propositional logic.

# First Order Logic

---

Constants	$Gold, Wumpus, [1, 2], [3, 1], \text{etc.}$
Predicates	$Adjacent(), Smell(), Breeze(), At()$
Functions	$Result()$
Variables	$x, y, a, t, \dots$
Connectives	$\wedge \vee \neg \Rightarrow \Leftrightarrow$
Equality	$=$
Quantifiers	$\forall \exists$

# Sentences

---

Brothers are siblings

$$\forall x, y \text{ Brother}(x, y) \Rightarrow \text{Sibling}(x, y)$$

“Sibling” is symmetric

$$\forall x, y \text{ Sibling}(x, y) \Leftrightarrow \text{Sibling}(y, x)$$

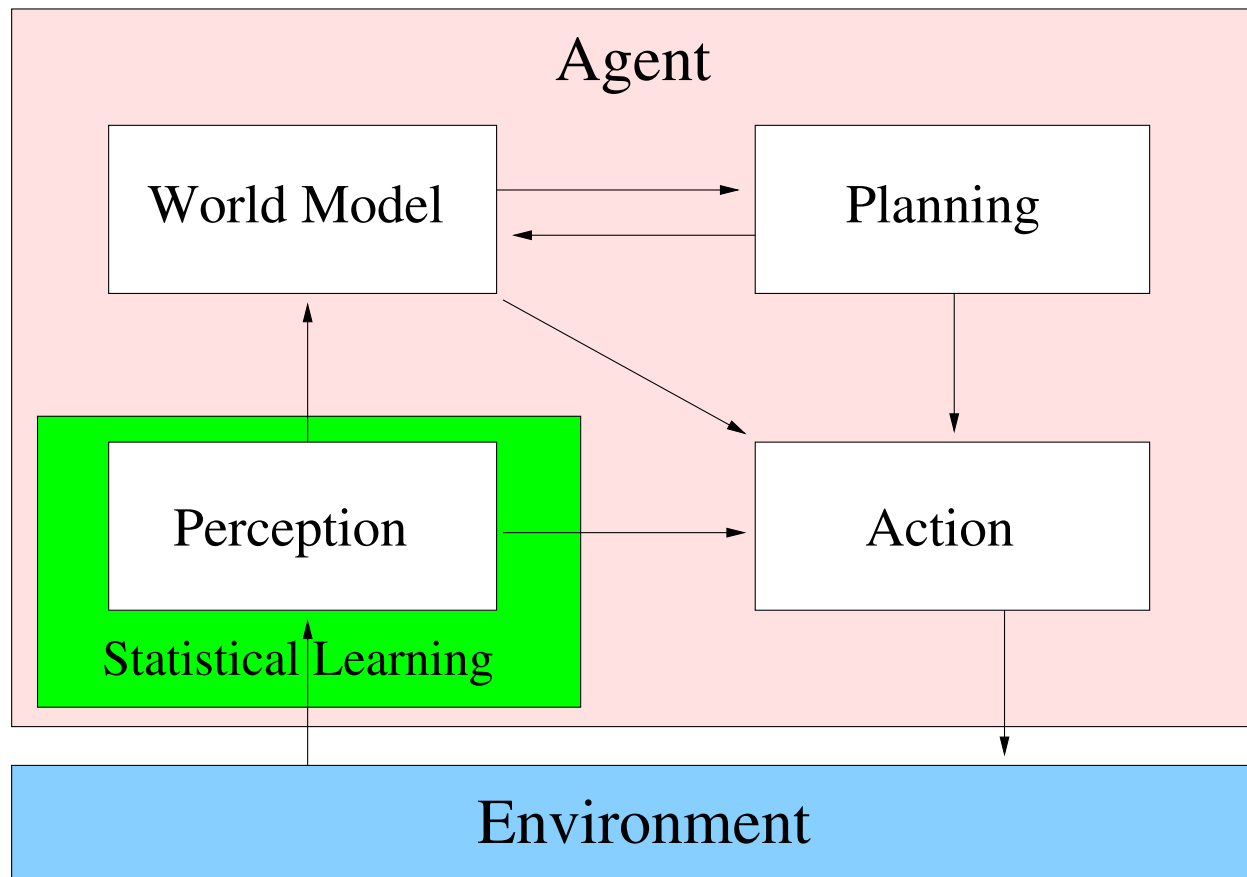
One's mother is one's female parent

$$\forall x, y \text{ Mother}(x, y) \Leftrightarrow (\text{Female}(x) \wedge \text{Parent}(x, y))$$

A first cousin is a child of a parent's sibling


$$\forall x, y \text{ FirstCousin}(x, y) \Leftrightarrow \exists p, ps \text{ Parent}(p, x) \wedge \text{Sibling}(ps, p) \wedge \text{Parent}(ps, y)$$

# Statistical Learning Agent



# Ockham's Razor

“The most likely hypothesis is the **simplest** one consistent with the data.”



inadequate


good compromise

over-fitting


Since there can be **noise** in the measurements, in practice need to make a tradeoff between simplicity of the hypothesis and how well it fits the data.



# Decision Tree



# Choosing an Attribute



Patrons is a “more informative” attribute than Type, because it splits the examples more nearly into sets that are “all positive” or “all negative”.

This notion of “informativeness” can be quantified using the mathematical concept of “entropy”.

A parsimonious tree can be built by minimizing the entropy at each step.

# Minimal Error Pruning

---

Should the children of this node be pruned or not?

Left child has class frequencies [2,4]

$$E = 1 - \frac{n+1}{N+k} = 1 - \frac{4+1}{6+2} = 0.375$$


Right child has  $E = 0.333$

Parent node has  $E = 0.444$

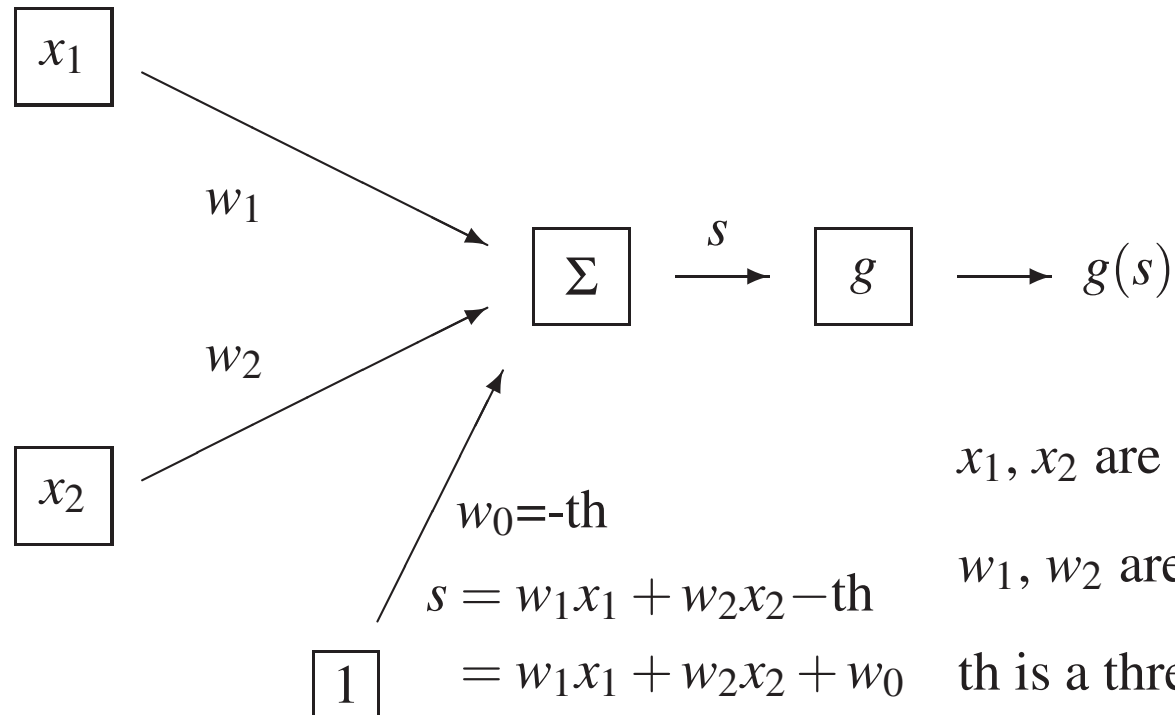
Average for Left and Right child is

$$E = \frac{6}{7}(0.375) + \frac{1}{6}(0.333) = 0.413$$

Since  $0.413 > 0.375$ , children should be pruned.



# Rosenblatt Perceptron



$x_1, x_2$  are inputs

$w_1, w_2$  are synaptic weights

th is a threshold

$w_0$  is a **bias** weight

$g$  is transfer function

# Perceptron Learning Rule

---

Adjust the weights as each input is presented.

recall:  $s = w_1x_1 + w_2x_2 + w_0$

if  $g(s) = 0$  but should be 1,

$$w_k \leftarrow w_k + \eta x_k$$

$$w_0 \leftarrow w_0 + \eta$$

$$\text{so } s \leftarrow s + \eta \left(1 + \sum_k x_k^2\right)$$

if  $g(s) = 1$  but should be 0,

$$w_k \leftarrow w_k - \eta x_k$$

$$w_0 \leftarrow w_0 - \eta$$


$$\text{so } s \leftarrow s - \eta \left(1 + \sum_k x_k^2\right)$$

otherwise, weights are unchanged. ( $\eta > 0$  is called the **learning rate**)

**Theorem:** This will eventually learn to classify the data correctly, as long as they are **linearly separable**.

# Multi-Layer Neural Networks

---



# Gradient Descent

---

We define an **error function**  $E$  to be (half) the sum over all input patterns of the square of the difference between actual output and desired output

$$E = \frac{1}{2} \sum (z - t)^2$$

If we think of  $E$  as height, it defines an error **landscape** on the weight space. The aim is to find a set of weights for which  $E$  is very low. This is done by moving in the steepest downhill direction.

$$w \leftarrow w - \eta \frac{\partial E}{\partial w}$$

Parameter  $\eta$  is called the **learning rate**.

# Probability and Uncertainty

	<i>toothache</i>		$\neg$ <i>toothache</i>	
	<i>catch</i>	$\neg$ <i>catch</i>	<i>catch</i>	$\neg$ <i>catch</i>
<i>cavity</i>	<b>.108</b>	<b>.012</b>	<b>.072</b>	<b>.008</b>
$\neg$ <i>cavity</i>	<b>.016</b>	<b>.064</b>	<b>.144</b>	<b>.576</b>

$$\begin{aligned}
 P(\neg \text{cavity} | \text{toothache}) &= \frac{P(\neg \text{cavity} \wedge \text{toothache})}{P(\text{toothache})} \\
 &= \frac{0.016 + 0.064}{0.108 + 0.012 + 0.016 + 0.064} = 0.4
 \end{aligned}$$



# Bayes' Rule

---

Product rule  $P(a \wedge b) = P(a|b)P(b) = P(b|a)P(a)$

$$\rightarrow \text{Bayes' rule } P(a|b) = \frac{P(b|a)P(a)}{P(b)}$$

Useful for assessing **diagnostic** probability from **causal** probability:

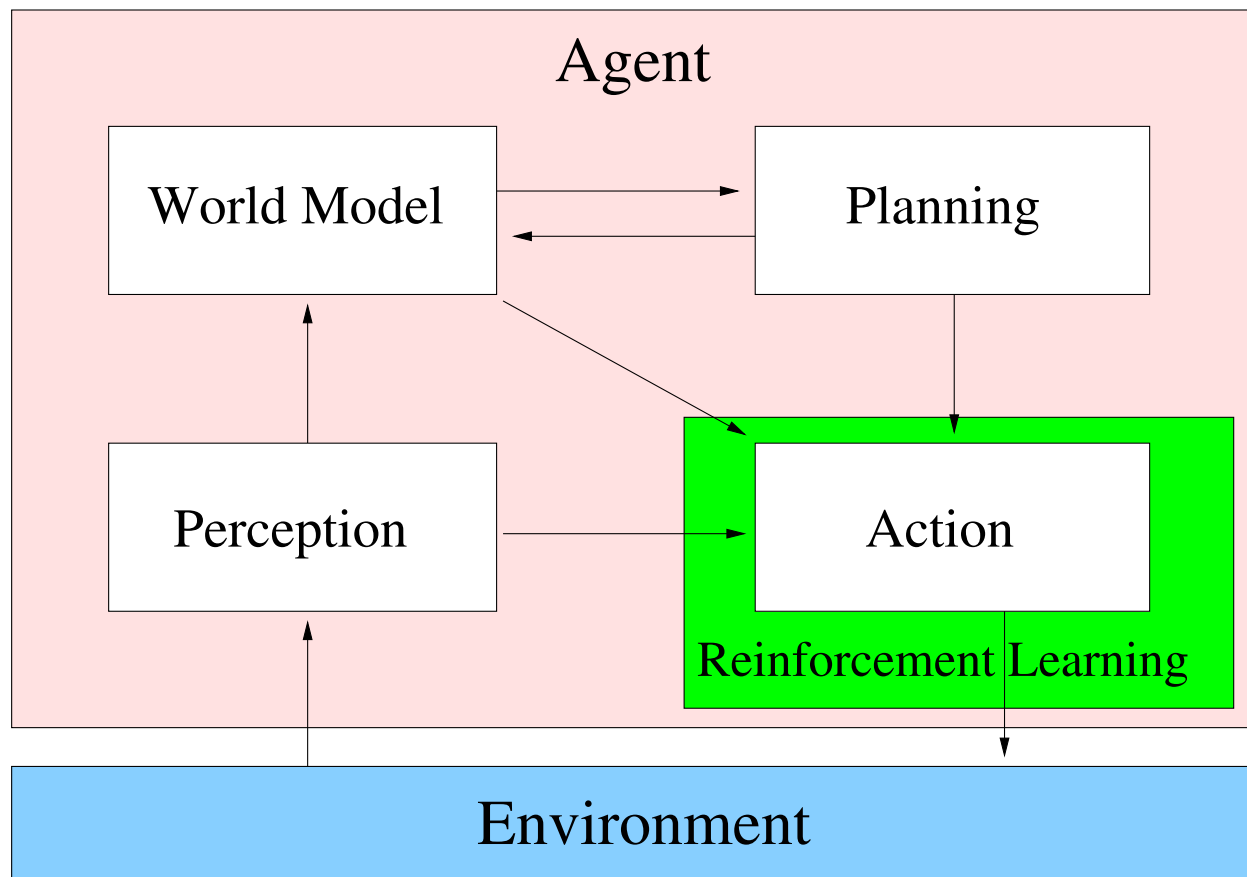
$$P(\text{Cause}|\text{Effect}) = \frac{P(\text{Effect}|\text{Cause})P(\text{Cause})}{P(\text{Effect})}$$

e.g., let  $M$  be meningitis,  $S$  be stiff neck:

$$P(m|s) = \frac{P(s|m)P(m)}{P(s)} = \frac{0.8 \times 0.0001}{0.1} = 0.0008$$

Note: posterior probability of meningitis still very small!

# Reinforcement Learning Agent



# Q-Learning

---

For each  $s \in S$ , let  $V^*(s)$  be the maximum discounted reward obtainable from  $s$ , and let  $Q(s, a)$  be the discounted reward available by first doing action  $a$  and then acting optimally.

Then the optimal policy is

$$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$$

where

$$Q(s, a) = r(s, a) + \gamma V^*(\delta(s, a))$$

then

$$V^*(s) = \max_a Q(s, a),$$

so

$$Q(s, a) = r(s, a) + \gamma \max_b Q(\delta(s, a), b)$$

which allows us to iteratively approximate  $Q$  by

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_b \hat{Q}(\delta(s, a), b)$$

# Final Exam

---

- similar to Tutorial Questions, but in Multiple Choice format
- Questions 1-24 one mark each, Questions 25-40 two marks each
- for each question, choose the ONE BEST Answer
- no marks taken off for wrong answers
- all modules covered in roughly equal proportion
- quick questions are at the beginning; longer questions at the end

## Sample 1-mark Question

---

Completeness of a search algorithm answers the question:

- (a) Is the algorithm guaranteed to find a solution when there is one?
- (b) Does the strategy find the solution that has the lowest path cost of all solutions?
- (c) How long does it take to find a solution?
- (d) How much memory is needed to perform the search?

## Sample 2-mark Question

---

Consider this joint probability distribution:

	short		$\neg$ short	
	wide	$\neg$ wide	wide	$\neg$ wide
striped	0.07	0.05	0.08	0.12
$\neg$ striped	0.14	0.17	0.12	0.25

Compute (to two decimal places):  $\text{Prob}(\text{short} \vee \neg \text{wide} \mid \text{striped})$

- (a) 0.50
- (b) 0.63
- (c) 0.75
- (d) 0.80

# Beyond COMP9414/3411

---

- COMP9444 Neural Networks and Deep Learning
- COMP9417 Machine Learning and Data Mining
- COMP4418 Knowledge Representation and Reasoning
- COMP3431 Robotic Software Architecture
- COMP9517 Machine Vision
- 4th Year Thesis topics

# UNSW myExperience Survey

---

Please remember to fill in the UNSW myExperience Survey.



# COMP3411/9414/9814 Artificial Intelligence

---

QUESTIONS?

# COMP3411/9414/9814 Artificial Intelligence

---

GOOD LUCK!